

# starter-code

April 1, 2024

## 0.1 Exploration

```
[256]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df_train = pd.read_csv("./data/train.csv")
df_test = pd.read_csv("./data/test.csv")

df_combined = pd.concat([df_train, df_test], axis=0)

df = df_combined

df_fraud = df[df.is_fraud==1]

print('Shape for preprocessed train dataset: \n', df.shape)
print('Shape for only-fraud train dataset: \n', df_fraud.shape)

fig = plt.figure(figsize=(10,6))
# sns.set()
colors = ["#C04C36", "#00163E"]
custom_params = {"axes.spines.right": False, "axes.spines.top": False}
sns.set_theme(style="ticks", rc=custom_params)
sns.set_context("notebook", font_scale=1.25)
sns.set_palette(sns.color_palette(colors))

# Create count plot with region on the y-axis
g = sns.countplot(y = 'city',
                  data=df_fraud,
                  hue='gender',
                  width=0.8,
                  order=df_fraud.city.value_counts(sort=True, ascending=False).
                    ↪head(10).index)

# Set title, label, legend
g.set_title('Top 10 cities in Number of fraud vs Gender', fontdict = {
    ↪'fontsize': 16, 'fontweight': 'bold'})
```

```

g.set_xlabel('Count', fontsize=15, fontweight='bold')
g.set_ylabel('City', fontsize=15, fontweight='bold')
g.legend(prop={'weight': 'bold'})

# Show plot
plt.show()

```

Shape for preprocessed train dataset:

(625184, 23)

Shape for only-fraud train dataset:

(1877, 23)



## 0.2 Feature Extraction

```

[311]: import pandas as pd
from sklearn.preprocessing import LabelEncoder
from geopy.distance import great_circle
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

def calculate_age(dob, trans_date):
    age = trans_date.year - dob.year - ((trans_date.month, trans_date.day) <
    ↪ (dob.month, dob.day))
    return age

```

```

def calculate_distance(row):
    customer_loc = (row['lat'], row['long'])
    merchant_loc = (row['merch_lat'], row['merch_long'])
    return great_circle(customer_loc, merchant_loc).km

def cluster_locations(df):
    print("Performing location clustering...")
    # Scaling the location data
    scaler = StandardScaler()
    loc_df_scaled = scaler.fit_transform(df[['lat', 'long']])
    kmeans = KMeans(n_clusters=20, random_state=0, n_init='auto').
    fit(loc_df_scaled)
    df['location_cluster'] = kmeans.labels_
    print("Location clustering completed.")

    return df

def calculate_rfm_features(df, window_days):
    """Calculates frequency and average monetary value within a time window"""

    now = df['trans_date_trans_time'].max()
    window_start = now - pd.Timedelta(days=window_days)

    window_transactions = df[df['trans_date_trans_time'] >= window_start]

    rfm_features = window_transactions.groupby('cc_num')[['amt']].agg(['count',
    ↪ 'mean'])
    rfm_features.columns = [f'freq_{window_days}days',
    ↪ f'avg_amt_{window_days}days']

    return rfm_features

def calculate_merchant_risk_scores(df, window_days, delay_days):
    """Calculates terminal risk scores within time windows, with a delay"""

    now = df['trans_date_trans_time'].max()
    window_end = now - pd.Timedelta(days=delay_days)
    window_start = window_end - pd.Timedelta(days=window_days)

    window_transactions = df[(df['trans_date_trans_time'] >= window_start) &
    ↪ (df['trans_date_trans_time'] < window_end)]

    terminal_risk = window_transactions.groupby('merchant')['is_fraud'].mean()
    terminal_risk.name = f'risk_score_{window_days}days'

    return terminal_risk

```

```

def calculate_category_risk_factors(df, window_days, delay_days):
    """Calculates category risk factors within time windows, with a delay"""

    now = df['trans_date_trans_time'].max()
    window_end = now - pd.Timedelta(days=delay_days)
    window_start = window_end - pd.Timedelta(days=window_days)

    window_transactions = df[(df['trans_date_trans_time'] >= window_start) &
    ↪(df['trans_date_trans_time'] < window_end)]

    category_risk = window_transactions.groupby('category')['is_fraud'].mean()
    category_risk.name = f'category_risk_{window_days}days'

    return category_risk

def extract_time_category_features(df):
    # Convert transaction datetime to hour of the day
    df['trans_hour'] = df['trans_time'].apply(lambda x: x.hour)

    df['category_hour'] = df['category'].astype(str) + "_" + df['trans_hour'].
    ↪astype(str)

    # To use these new interaction features in a model, you need to encode them
    # This encoding could be label encoding or one-hot encoding. Here's an
    ↪example with label encoding:

    label_encoder = LabelEncoder()
    df['category_hour_encoded'] = label_encoder.
    ↪fit_transform(df['category_hour'])

    return df

def calculate_odds_ratios(df, category_col='category_encoded'):
    # Compute the number of fraud and non-fraud transactions
    fraud_counts = df[df['is_fraud'] == 1][category_col].value_counts().
    ↪sort_index()
    non_fraud_counts = df[df['is_fraud'] == 0][category_col].value_counts().
    ↪sort_index()

    # Calculate the odds for each category
    odds = fraud_counts / non_fraud_counts

    # Calculate the overall odds of fraud
    overall_odds = df['is_fraud'].mean()

    # Calculate odds ratios (odds for each category / overall odds)

```

```

odds_ratios = odds / overall_odds

# Replace infinite values with a large number
odds_ratios = odds_ratios.replace(np.inf, np.finfo(float).max)

return odds_ratios

def process(df):
    print("Starting data processing...")
    # Convert to datetime just once
    df['dob'] = pd.to_datetime(df['dob'], format='%d/%m/%Y')
    df['trans_date_trans_time'] = pd.to_datetime(df['trans_date_trans_time'],
    ↪format='%d/%m/%Y %H:%M')

    # Extracting date and time for each row
    df['trans_date'] = df['trans_date_trans_time'].dt.date
    df['trans_time'] = df['trans_date_trans_time'].dt.time

    print("Cleaning data and calculating basic statistics...")

    # Transaction amount statistics for each cardholder
    agg_funcs = ['mean', 'std', 'min', 'max']
    cardholder_stats = df.groupby('cc_num')['amt'].agg(agg_funcs).
    ↪rename(columns=dict(zip(agg_funcs, ['mean', 'std', 'min', 'max'])))
    df = df.join(cardholder_stats, on='cc_num', rsuffix='_cardholder')

    # Encode categorical features
    categorical_features = ['merchant', 'city', 'state', 'job']
    df[categorical_features] = df[categorical_features].apply(LabelEncoder().
    ↪fit_transform)

    df['category_encoded'] = LabelEncoder().fit_transform(df['category'])
    df['gender_encoded'] = LabelEncoder().fit_transform(df['gender'])

    # Calculate odds ratios for each category
    odds_ratios = calculate_odds_ratios(df, category_col='category_encoded')
    df['category_odds_ratio'] = df['category_encoded'].map(odds_ratios).
    ↪fillna(1)

    print("Calculating age at transaction and distances...")
    # Age at transaction
    df['age_at_transaction'] = df.apply(lambda x: calculate_age(x['dob'],
    ↪x['trans_date']), axis=1)

    # Geographical feature: Distance between customer and merchant
    df['cust_merch_distance'] = df.apply(calculate_distance, axis=1)

```

```

df['trans_weekend'] = df['trans_date_trans_time'].dt.weekday >= 5
df['trans_night'] = df['trans_date_trans_time'].dt.hour.apply(lambda x: 1
↳ if 0 <= x < 6 or 20 <= x < 24 else 0)

for window_days in [1, 7, 30]:
    rfm_features = calculate_rfm_features(df.copy(), window_days)
    df = df.merge(rfm_features, how='left', on='cc_num')

    terminal_risk = calculate_merchant_risk_scores(df.copy(), window_days,
↳ delay_days=7)
    df = df.merge(terminal_risk, how='left', on='merchant')

    category_risk = calculate_category_risk_factors(df.copy(), window_days,
↳ delay_days=7)
    df = df.merge(category_risk, how='left', on='category')

    # Perform location clustering
    df = cluster_locations(df)

    df = extract_time_category_features(df)

    df.drop(columns=['first', 'last', 'street', 'gender',
↳ 'trans_date_trans_time', 'trans_date', 'trans_time', 'cc_num'], inplace=True)

    print("Data processing completed.")

    return df

# Load the dataset
trainingSet = pd.read_csv("./data/train.csv")

# Process the DataFrame
print("Processing training set...")
train_processed = process(trainingSet)
print("Training set processed.")

# Load test set
submissionSet = pd.read_csv("./data/test.csv")

# Merge on Id so that the test set can have feature columns as well
testX = pd.merge(train_processed, submissionSet, left_on='Id', right_on='Id')
testX = testX.drop(columns=['is_fraud_x'])
testX = testX.rename(columns={'is_fraud_y': 'is_fraud'})

# The training set is where the score is not null
trainX = train_processed[train_processed['is_fraud'].notnull()]

```

```
# Save the datasets with the new features for easy access later
testX.to_csv("./data/X_test.csv", index=False)
trainX.to_csv("./data/X_train.csv", index=False)
```

Processing training set...  
 Starting data processing...  
 Cleaning data and calculating basic statistics...  
 Calculating age at transaction and distances...  
 Performing location clustering...  
 Location clustering completed.  
 Data processing completed.  
 Training set processed.

### 0.3 Creating your model

```
[312]: import pickle
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, f1_score
from xgboost import XGBClassifier

# Load training set with new features into DataFrame
X_train = pd.read_csv("./data/X_train.csv")

# Split training set into training and testing set
X_train, X_test, Y_train, Y_test = train_test_split(
    X_train.drop(['is_fraud', 'Id'], axis=1),
    X_train['is_fraud'],
    test_size=0.1,
    random_state=42
)

# This is where you can do more feature selection
X_train_processed = X_train._get_numeric_data()
print(X_train_processed.columns)
X_test_processed = X_test._get_numeric_data()

# Define XGBoost model (you might want to tune these)
xgb_model = XGBClassifier(random_state=0, use_label_encoder=False,
    ↪eval_metric='logloss')

# Learn the model
xgb_model.fit(X_train_processed, Y_train)
```

```

# Pickle model
with open('xgboost_model.obj', 'wb') as f:
    pickle.dump(xgb_model, f)

# Evaluate on the testing set
Y_test_predictions = xgb_model.predict(X_test_processed)
print("Accuracy on testing set = ", accuracy_score(Y_test, Y_test_predictions))
print("F1 score on testing set = ", f1_score(Y_test, Y_test_predictions))

# Plot a confusion matrix
cm = confusion_matrix(Y_test, Y_test_predictions)
sns.heatmap(cm, annot=True)
plt.title('Confusion matrix of the classifier')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

```

```

Index(['merchant', 'amt', 'city', 'state', 'zip', 'lat', 'long', 'city_pop',
      'job', 'unix_time', 'merch_lat', 'merch_long', 'mean', 'std', 'min',
      'max', 'category_encoded', 'gender_encoded', 'category_odds_ratio',
      'gender_odds_ratio', 'age_at_transaction', 'cust_merch_distance',
      'trans_weekend', 'trans_night', 'freq_1days', 'avg_amt_1days',
      'risk_score_1days', 'category_risk_1days', 'freq_7days',
      'avg_amt_7days', 'risk_score_7days', 'category_risk_7days',
      'freq_30days', 'avg_amt_30days', 'risk_score_30days',
      'category_risk_30days', 'location_cluster', 'trans_hour',
      'category_hour_encoded'],
      dtype='object')

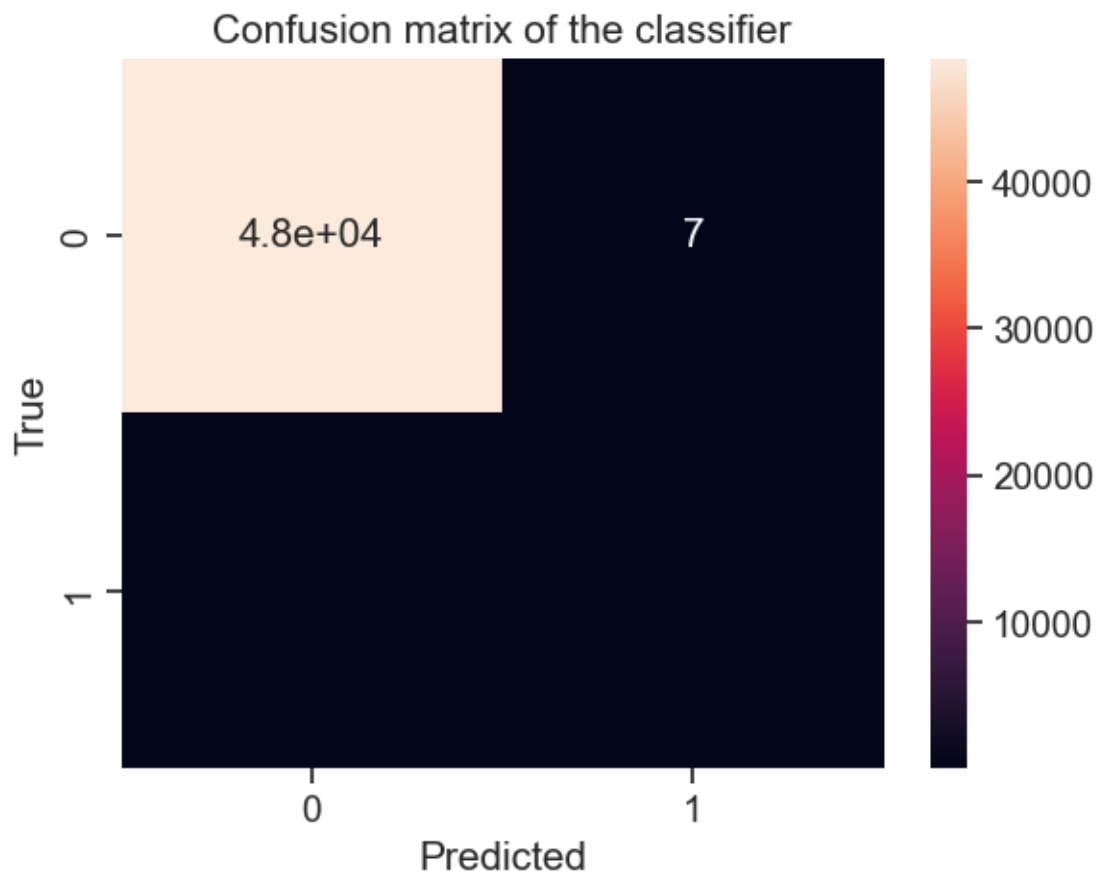
```

`use\_label\_encoder` is deprecated in 1.7.0.

Accuracy on testing set = 0.9993830461070209

F1 score on testing set = 0.9019607843137255





```
[308]: X_train_processed.head()
```

```
[308]:
```

	merchant	amt	city	state	zip	lat	long	city_pop	job	\
230198	70	37.81	690	47	54559	46.4959	-90.4383	795	446	
201736	271	16.46	542	2	71960	34.4596	-93.6743	1383	265	
426780	100	103.05	733	22	55080	45.6675	-93.2433	2607	254	
17706	178	94.31	243	30	7022	40.8170	-74.0000	13835	354	
180397	33	133.62	123	47	53924	43.4987	-90.2796	1360	302	

	unix_time	...	avg_amt_7days	risk_score_7days	category_risk_7days	\
230198	1375853554	...	85.199149	0.0	0.000000	
201736	1387011410	...	71.865179	0.0	0.001496	
426780	1375130700	...	45.662535	0.0	0.001482	
17706	1377325280	...	60.025455	0.0	0.000000	
180397	1378635449	...	87.870625	0.0	0.001948	

	freq_30days	avg_amt_30days	risk_score_30days	category_risk_30days	\
230198	207.0	93.963188	0.010526	0.001185	
201736	202.0	131.916040	0.007353	0.006904	

426780	236.0	72.610085	0.000000	0.001551
17706	161.0	55.785217	0.000000	0.000639
180397	159.0	98.946667	0.005464	0.003979

	location_cluster	trans_hour	category_hour_encoded
230198	0	5	175
201736	11	8	154
426780	0	20	116
17706	6	6	20
180397	9	10	218

[5 rows x 40 columns]

```
[313]: # After fitting your XGBoost model (xgb_model)
importances = xgb_model.feature_importances_
feature_names = X_train_processed.columns # Replace with the names of your
↳ features

sorted_features = [feature for _, feature in sorted(zip(importances,
↳ feature_names), reverse=True)]
for i in sorted_features:
    print(i)

import matplotlib.pyplot as plt
import xgboost as xgb


# If you use 'gain' as your importance_type:
xgb.plot_importance(xgb_model, importance_type='gain', max_num_features=10)
plt.show()
```

```
category_encoded
trans_night
category_odds_ratio
trans_hour
amt
category_risk_7days
min
mean
age_at_transaction
std
freq_30days
category_risk_30days
gender_encoded
max
risk_score_30days
city_pop
freq_1days
```

```

category_hour_encoded
unix_time
avg_amt_30days
freq_7days
avg_amt_7days
avg_amt_1days
risk_score_7days
lat
state
trans_weekend
long
job
city
location_cluster
zip
merch_long
cust_merch_distance
merch_lat
merchant
risk_score_1days
gender_odds_ratio
category_risk_1days

```



```

[280]: import shap

# Fit an explainer
explainer = shap.TreeExplainer(xgb_model)

```

```
test_IDs = df['Id'].copy()

X_shap = X_train_processed.copy()

shap_values = explainer.shap_values(X_shap)


print(shap_values.shape)
print(X_train_processed.shape)

shap.summary_plot(shap_values, X_train_processed)
```

(437628, 36)

(437628, 36)

No data for colormapping provided via 'c'. Parameters 'vmin', 'vmax' will be ignored



#### 0.4 Create the Kaggle submission

```
[303]: X_submission = pd.read_csv("./data/X_test.csv")
test_IDs = X_submission['Id']
X_submission = X_submission.drop(columns=['is_fraud', 'Id'])
X_submission_processed = X_submission._get_numeric_data()
```

```

X_submission['is_fraud'] = xgb_model.predict(X_submission_processed)
X_submission.is_fraud = X_submission.is_fraud.astype(int)
X_submission['Id'] = test_IDs
submission = X_submission[['Id', 'is_fraud']]
submission.to_csv("./data/submission.csv", index=False)

```

Now you can upload the submission.csv to kaggle

```

[291]: import xgboost as xgb
import matplotlib.pyplot as plt

# Load the model
# Replace with the correct path or code to load your model
xgb_model = pickle.load(open('XGBoost_best_model.pkl', 'rb'))

# Plot feature importances
xgb.plot_importance(xgb_model, max_num_features=10) # Adjust to display the
↳ number of top features you want to see
plt.show()

# If your categories are one-hot encoded, their feature importances will be
↳ scattered across multiple features,
# and you might need to sum these importances to get the overall importance for
↳ the category feature.

```

